

Linear Time Generation of Simulated Wireless Sensor Networks with Random Geometric Graphs

Luke Wood

April 10, 2018

1 Executive Summary

Wireless Sensor Networks (WSNs) are a group of ad hoc wireless devices that communicate amongst themselves. WSNs have tons of applications coming in the near future ranging from traffic sensing networks to weather reading sensors on the surface of Mars[3]. They are incredibly expensive to develop and test which makes them a great candidate for simulation as a preliminary way to gather data. Vlady Ravelomananana and Hichem Kenniche from the University of Paris first explored the concept of using random geometric graphs (RGGs) to attempt to model wireless sensor networks [2]. At the moment random geographic graphs are the state of the art method for simulating WSNs. In this project I will use RGGs as a way to gather valuable information about how wireless sensor networks will possibly function and communicate.

1.1 Introduction and Summary

Through a series of three reports I will be analyzing RGGs in the following ways as an attempt to gain some insight into the behavior of wireless sensor networks:

1. Generating RGGs on the geometries of a unit square, unit disc, and unit sphere
2. Color the generated graph in linear time using smallest vertex last ordering[4]
3. Find the terminal clique in the generated RGGs
4. Find a selection of bipartite subgraphs produced by an algorithm for coloring

This report is the first of the series and will describe an implementation for a linear time algorithm to generate graphs consisting of N vertices with an average degree of A on the geometric topologies of: unit square, unit disc, unit sphere. Some tables are included to display some metrics for generated RGGs of various size to display the performance of the generation algorithm. The square

Benchmark	N	A	R	Topology	Avg. Degree	Min Deg	Max Degree	When Removed
1	1000	32	0.100925	Square	29	6	52	22
2	8000	64	0.050463	Square	61	13	95	39
3	16000	32	0.025231	Square	31	7	55	22
4	64000	64	0.017841	Square	63	15	96	44
5	64000	128	0.025231	Square	125	31	172	73
6	128000	64	0.012616	Square	63	11	103	44
7	128000	128	0.017841	Square	126	31	172	74
8	8000	64	0.044721	Disc	61	24	94	44
9	64000	64	0.015811	Disc	63	20	97	42
10	64000	128	0.022361	Disc	125	39	166	72
11	16000	64	0.063246	Sphere	63	40	97	44
12	32000	128	0.063246	Sphere	128	84	176	85
13	64000	128	0.044721	Sphere	128	77	171	78

Table 1: Benchmark Data Generated by the Implementation

topology generation algorithm runs in approximately 15 seconds for graph size 100000 which is fast given the large input size.

The disc topology also runs in linear time but the constant factor is much larger than that of the square topology. This is due to some overhead that I incurred when generating the points themselves. Despite the algorithm being slower than the square version, it still terminated in 48 seconds for 100000 nodes. At the moment this is the worst generation algorithm and could be improved by using mathematical formulas instead of repeated guessing to generate points.

Sphere generation was also a linear time algorithm. Unlike the disc topology I generated the points using a mathematical formula instead of brute force finding points. This yielded a runtime of under 35 seconds for generating 32000 points.

One interesting thing to notice in the tables is that as the number of nodes grows the real average degree converges on the expected average degree. The lower the number of nodes the further from the expected average degree the real number is. This happens due to there not being enough nodes in the graph for the real radius to reach the expected value. The random error in the point generation is more apparent when there are fewer nodes in the graph.

Reference table 1.1 for the full benchmarks.

1.2 Programming Environment Description

The implementation of the algorithm used to gather the data supporting this report was gathered on a 15 inch Macbook pro 2017 with a 2.9 GHz Intel Core i7 processor and 16 GB of RAM. The computer is running macOS High Sierra. The graph generation is written in python 3 as generating and connection a graph is not super computationally expensive with even decently large inputs such as 100000 (assuming $O(n)algorithms$). The later algorithms may be implemented

in a different language such as Elixir to get high levels of concurrency and higher efficiency due to type inference (as opposed to python's dynamic typing).

2 Reduction to Practice

This section will describe the transition from theory to implementation. This section will also give a detailed analysis of the algorithms used in this project as well as their asymptotic runtimes.

2.1 Data Structure Design

The generation portion of this project uses several different data structures. The first one is a python object of custom type node. This serves as a tuple of values consisting of a list of dimensions, a list of nodes, and a node number. The first two are used during graph generation and the latter during conversion to an adjacency list. This data structure could be used interchangeably with a statically indexed tuple instead of an object to avoid any overhead associated with objects in python, however the readability gained from using a custom node class heavily outweighs the marginal performance benefit gained from using a statically indexed tuple. Both of these data structures provide $O(1)$ read and write operations. The python wiki states that dictionary access time is amortized $O(1)$, however in the worst case can be $O(n)$ [5]. That being said, most interpreters will convert objects with known attributes into statically indexed arrays so it is likely that there would be negligible performance difference anyways. If the goal were to generate gigantic graphs then an argument could be made to switch over the statically indexed tuples to reduce the access time for attributes by a bit as tuples use static memory address offsets.

The second mentioned data structure is the adjacency list. Adjacency lists are an efficient graph representation that we will use in the subsequent reports. Adjacency lists require only $O(v * e)$ storage as opposed to the $O(v^2)$ required for adjacency matrixes. This is handy in situations where the expected average number of edges is significantly lower than the number of nodes in the graph. Despite the huge potential savings on storage, the only sacrifice adjacency lists make is in the lookup operation to determine if there is an edge between two nodes. This takes $O(e)$ in an adjacency list as opposed to $O(1)$ in an adjacency matrix.

2.2 Algorithm Description

This section gives a detailed analysis of the algorithms used in the graph generation and coloring.

Square Point Generation The algorithm to generate the points in the unit square topology is simple. The steps are as follows:

1. $x=\text{random}(0, 1)$, $y=\text{random}(0, 1)$

2. add (x,y) to a list of points
3. repeat steps 1 and 2 until N points are created

That is all that is necessary to get a random point in the unit square. This algorithm is $\Theta(n)$.

Disc Point Generation Generating the points on the disc topology is slightly more challenging than generating the points for the square topology. The algorithm for generating points used in the implementation described in this point is as follows:

1. $x = \text{random}(0, 1)$, $y = \text{random}(0, 1)$
2. if the distance from (x,y) to (.5, .5) is $\leq .5$, add (x,y) to a list of points
3. repeat steps 1 and 2 until N points are created

The issue with this algorithm is the nondeterminism involved in creating each point. When generating large numbers of points lots of work is wasted on generating the unused random points. In future iterations of this project a mathematical approach may be used to generate these points. Despite the non-determinism, this algorithm is still amortized $\Theta(n)$.

A possible approach to fixing this implementation would be to pick a radius between 0 and 1, and then rotate by a random number of radians between 0 and 2π . This was not implemented as I was unsure if this would produce a uniform distribution.

Sphere Point Generation Generating points on the surface of the unit sphere can be solved using math. The steps for the following algorithm are as follows:

1. $u = (\text{random}() * 2) - 1$, $\theta = \text{random}() * 2 * \pi$
2. $x = \sqrt{1 - u^2} * \cos(\theta)$, $y = \sqrt{1 - u^2} * \sin(\theta)$, $z = u$
3. append (x, y, z) to a list of points.
4. repeat steps 1-3 N times

In the implementation associated with this report the values x, y, and z are transformed to form a sphere around (.5, .5, .5) as opposed to (0, 0, 0) as it makes the connection algorithm simpler. This algorithm is $\Theta(n)$.

Topology	Radius Equation
Unit Square	$r = \sqrt{\frac{d(G)}{N\pi}}$
Unit Disc	$r = \sqrt{\frac{d(G)}{N}}$
Unit Sphere	$r = \sqrt{\frac{2 * d(G)}{N}}$

2.3 Node Connection

In order to ensure that average degree of the nodes is close to the desired average degree, a radius of a specific length can be defined surrounding each node. The nodes within that radius of each other are then connected in the graph. The formulas to find the radius for each topology is derived from the equations found in the paper Bipartite Grid Partitioning of a Random Geometric Graph[1]. The formula used to find this radius varies for each graph topology and can be found in the table 2.3.

The brute force algorithm for generating and connecting RGGs in this way is $O(n^2)$. This quickly became problematic as the algorithm took upwards of 200 seconds to run on input size of only 12,000. Figure 2.3 shows this issue. Instead, the implementation uses a bucket method to narrow down the required number of comparisons. The idea for this algorithm comes from the paper Bipartite Grid Partitioning of a Random Geometric Graph[1]. The algorithm can be broken down into steps as follows:

1. create $(\sqrt{1/R} - 1)^2$ buckets
2. place buckets in a 2D grid arrangement
3. for each node, place the node in bucket (x, y) such that $x = \text{floor}(x * \text{num_buckets})$, $y = \text{floor}(y * \text{num_buckets})$
4. for each x, y such that $x = 0.. \text{buckets}$, $y = 0.. \text{buckets}$
5. check to see if each node in bucket (x, y) is within radius R to all others in bucket (x, y) as well as $(x+1, y-1)$, $(x+1, y)$, $(x, y+1)$, $(x+1, y+1)$
6. if two nodes are under radius R from each other, place an edge between them.

This algorithm is $O(n)$ given a small enough radius and a sufficient number of buckets. This means that with sufficiently many nodes and as long as the expected average degree remains low the algorithm will be linear. The runtimes for the implementation of this algorithm can be found in figure 2.3.

Shortest Last Vertex Ordering In order to color the graph in linear time, the shortest last vertex ordering is used as a heuristic to get a reasonable coloring




Figure 1: Runtimes of the $O(n^2)$ Algorithm




Figure 2: Runtimes of the $O(n)$ Algorithm

of the graph. To compute the ordering in $O(V+E)$ time, the following algorithm is used. Graphics are attached to visualize each step of the process.

To begin the graph is uncolored, this can be seen in figure 2.3. The next step is to initialize the ordering of the nodes. To do so, the following data structures are created:

1. A map mapping from degrees to nodes which are of that degree - this should be a data structure with constant time access.
2. Each item mapped to in the previous map should be a set to allow for constant time access.
3. A map mapping from node numbers to the degree of that node - this should also be a data structure with constant time access.

Then, the following steps are executed for N iterations:

1. Select a node from a non-empty set mapped to by the smallest possible degree
2. Add this node to the ordering
3. Remove the node from the graph, this entails decrementing the degree of all neighbors and removing the node from the sets above.
4. Move all of the neighboring nodes one hash set lower in the map mapping from degree to nodes

Figure 2.3 shows this process for a graph with 20 nodes. After the ordering is complete, the ordering is then used to color the graph.

The coloring process is rather simple. Each node is taken from the ordering and is given the smallest color not yet assigned to any of it's neighbors. This is repeated until all nodes are colored. Figure 2.3 shows this process for the same example above. After this is complete, the graph will be fully colored. Figure 2.3 shows a fully graphed graph.

2.4 Algorithm Engineering

As mentioned above, the brute force algorithm to find adjacent nodes becomes slow at a fast rate. To fix this, the implementation presented uses linear time algorithms for all of it's operations. Table 2.4 compares the runtimes of the linear implementation to that of the quadratic implementation.

The $O(n)$ algorithm is far superior even on small input sizes such as 1000. The final implementation uses $O(n)$ time as well as $O(n * A)$ space where A=aveage edge density. This is accomplished by ensuring that each node is stored in a representation that only references adjacent edges. The implementation also only compares nodes to nodes in adjacent buckets to determine if an edge should be placed between them. Each node only has at most $5nr^2$ comparisons[1]. With reasonably spare graphs this maintains a runtime of $O(n)$ however with large enough radii this can begin to behave like $O(n^2)$. This yields an $O(n)$ time and $O(n)$ space on the type of inputs this implementation targets.

UnColored Unit Square Graph with N=20, R=.4




Figure 3: An Uncolored Graph on the square topology of N=20, and R=.4

Nodes	$O(n^2)$	$O(n)$
1000	0.258400	0.657174
2000	0.382116	2.630040
3000	0.473916	5.874501
5000	0.831753	16.809004
10000	1.560216	65.398991

Table 2: Runtimes of the $O(n)$ and $O(n^2)$ algorithms in seconds

Visualization of the removal of nodes during SLVO




Figure 4: An Example of Shortest Vertex Last Ordering

Coloring of Graph during SLVO Coloring




Figure 5: The Process of Coloring the Graph using SVLO

Colored Unit Square Graph with $N=20$, $R=.4$




Figure 6: A Fully Colored Graph

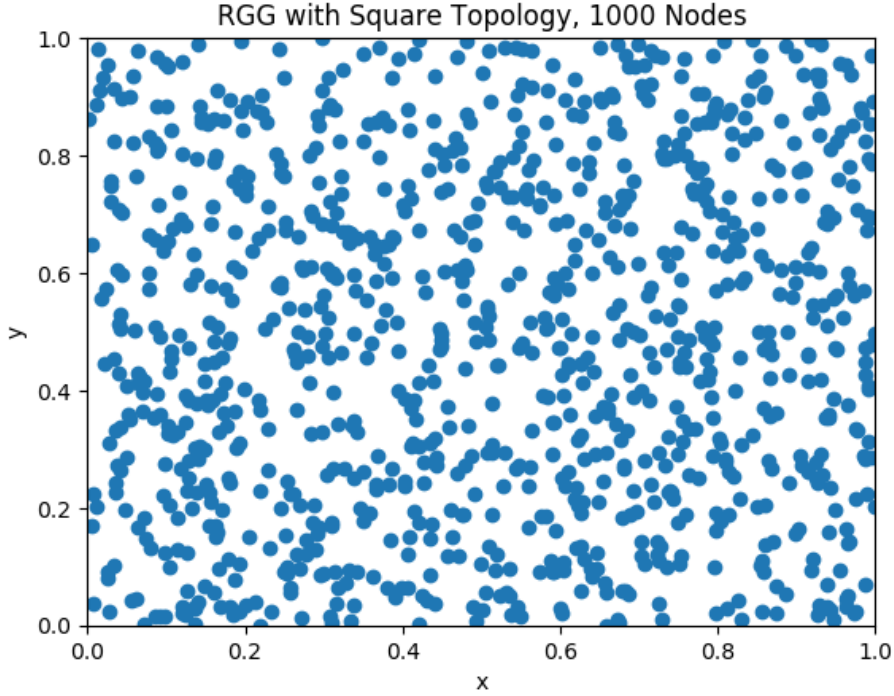


Figure 7: 1000 Points on the Square Topology

2.5 Verification

One way that we verified our results was checking the distribution of edge densities in our graph. We expect to see a gaussian distribution in the edge densities with the center being around our calculated radius. We can also verify the runtime of our algorithms by plotting the input size on the x axis and the runtime on the y axis. If we have a linear algorithm we should be able to fit the distribution of points to a linear equation with minimal error. Both of these verification methods were successful and can be seen in the results section of this report.

Visualizing the Points One way to validate that the points are distributing correctly is by plotting out the points in a scatter plot. Early on the implementation had a bug where the points distributed around the radius of the unit disc instead of evenly inside the unit disc. Using a scatter plot made this trivial to spot.

Figure 2.5 shows the distribution on the square topology, Figure 2.5 shows the distribution on the disc topology, and Figure 2.5 shows the distribution on the sphere topology.




Figure 8: 1000 Points on the Disc Topology




Figure 9: 1000 Points on the Sphere Topology




Figure 10: Edge Densities of an 8000 Node Graph with $E(\text{Degree})=64$ on Topology Square

Edge Density Checking the distribution of the edge densities is one way to verify that the edges are connecting as expected. The edge distributions should follow a gaussian distribution. Figures 2.5, 2.5, and 2.5 all show gaussian distributions for their respective topologies. All of the edge density distributions follow a gaussian distribution which is what is expected meaning that the edges likely connected the nodes in the expected way.

3 Result Summary

All of my algorithms ran in $O(n)$ time. There was variance between the runtime of the different topologies due to the way the points were generated combined with the calculation required to compute 2d distance vs 3d distance. Table 3 shows the runtime of all of the topologies in a table. Figures 3, 3, and 3 show the individual runtime charts. Figure 3 shows all of the runtimes of the different RGG generation algorithms overlaying each other. As discussed in the verification section, multiple plots have been used to verify the accuracy of the algorithms.




Figure 11: Edge Densities of an 8000 Node Graph with $E(\text{Degree})=64$ on Topology Disc

N	A	Square Runtime	Disc Runtime	Sphere Runtime
1000	64	0.444113	0.872433	2.983946
5000	64	1.811515	4.316646	6.846584
10000	64	2.931157	7.800921	10.893038
25000	64	7.639537	22.379140	30.186049
50000	64	16.041881	46.400750	63.725632
100000	64	30.260184	89.923685	125.474381

Table 3: Comparison of Runtimes of Generating the Different Topologies




Figure 12: Edge Densities of an 8000 Node Graph with $E(\text{Degree})=64$ on Topology Sphere




Figure 13: Runtimes of the $O(n)$ Square RGG Generation Algorithm




Figure 14: Runtimes of the $O(n)$ Disc RGG Generation Algorithm




Figure 15: Runtimes of the $O(n)$ Sphere RGG Generation Algorithm




Figure 16: Runtimes of the $O(n)$ RGG Generation Algorithms overlayed

References

- [1] Zizhen Chen and David W Matula. “Bipartite Grid Partitioning of a Random Geometric Graph”. In: *2017 13th International Conference on Distributed Computing in Sensor Systems (DCOSS)*. IEEE. 2017, pp. 163–169.
- [2] Hichem Kenniche and Vlado Ravelomananana. “Random geometric graphs as model of wireless sensor networks”. In: *Computer and Automation Engineering (ICCAE), 2010 The 2nd International Conference on*. Vol. 4. IEEE. 2010, pp. 103–107.
- [3] Dhia Mahjoub and David W Matula. “Employing $(1 - \varepsilon)$ dominating set partitions as backbones in wireless sensor networks”. In: *Proceedings of the Meeting on Algorithm Engineering & Experiments*. Society for Industrial and Applied Mathematics. 2010, pp. 98–112.
- [4] David W Matula and Leland L Beck. “Smallest-last ordering and clustering and graph coloring algorithms”. In: *Journal of the ACM (JACM)* 30.3 (1983), pp. 417–427.
- [5] PythonWiki. *TimeComplexity - Python Wiki*. 2017. URL: <https://wiki.python.org/moin/TimeComplexity> (visited on 06/05/2017).